

```
In [5]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import seaborn as sns
from sklearn import metrics


print("🏠 House Price Prediction Using Linear Regression Model from scikit-learn")
print("☕ Grab a cup of coffee and watch the magic happen!")


# Load the dataset
df = pd.read_csv('EcommerceCustomers.csv')
pd.set_option('display.max_columns', None)
print(df.head())
x = df[['Avg. Session Length', 'Time on App', 'Time on Website', 'Length of Membership']]
y = df['Yearly Amount Spent']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random_state=42)
lm = LinearRegression()
lm.fit(x_train, y_train)

# Evaluate the model
print("\n🔍 Model evaluation:")
print(f"🟢 Intercept (constant term): {lm.intercept_:.2f}")
print(f"🟢 Coefficients for each feature:")

# Create a DataFrame for coefficients
coeff_df = pd.DataFrame(lm.coef_, x.columns, columns=['Coefficient'])
print(coeff_df)
print("\n🔍 Interpretation of coefficients:")
for feature, coef in zip(x.columns, lm.coef_):
    print(f"💎 A one unit increase in '{feature}' is associated with an increase of {coef:.2f} in the yearly amount spent.")

predictions = lm.predict(x_test)
print("\n🔍 House price predictions:")
print(predictions)
```

 House Price Prediction Using Linear Regression Model from scikit-learn...

 Grab a cup of coffee and watch the magic happen!

Email \

```
0      mstephenson@fernandez.com
1      hduke@hotmail.com
2      pallen@yahoo.com
3      riverarebecca@gmail.com
4      mstephens@davidson-herman.com
```

Address

Avatar \

```
0      835 Frank Tunnel\nWrightmouth, MI 82180-9605      Violet
1      4547 Archer Common\nDiazchester, CA 06566-8576      DarkGreen
2      24645 Valerie Unions Suite 582\nCobbborough, D...      Bisque
3      1414 David Throughway\nPort Jason, OH 22070-1220      SaddleBrown
4      14023 Rodriguez Passage\nPort Jacobville, PR 3...      MediumAquaMarine
```

	Avg. Session Length	Time on App	Time on Website	Length of Membership \
0	34.497268	12.655651	39.577668	4.082621
1	31.926272	11.109461	37.268959	2.664034
2	33.000915	11.330278	37.110597	4.104543
3	34.305557	13.717514	36.721283	3.120179
4	33.330673	12.795189	37.536653	4.446308

	Yearly Amount Spent
0	587.951054
1	392.204933
2	487.547505
3	581.852344
4	599.406092

 Model evaluation:

 Intercept (constant term): -1047.98

 Coefficients for each feature:

	Coefficient
Avg. Session Length	25.912259
Time on App	38.508126
Time on Website	0.288716
Length of Membership	61.161610

 Interpretation of coefficients:

◆ A one unit increase in 'Avg. Session Length' is associated with an increase of \$25.91 in the house price.

◆ A one unit increase in 'Time on App' is associated with an increase of \$38.51 in the house price.

◆ A one unit increase in 'Time on Website' is associated with an increase of \$0.29 in the house price.

◆ A one unit increase in 'Length of Membership' is associated with an increase of \$61.16 in the house price.

 House price predictions:

```
[456.54286407 403.04038845 409.4733783 591.19661352 589.78507712
548.65253346 577.22540667 714.97215336 473.60765513 545.80428474
338.10305743 500.26666332 552.89535503 409.57836003 764.88121506
545.59508474 692.68128903 507.24854118 572.83232081 572.96763324
397.66463172 554.77893276 458.29897396 482.60971744 558.96461845
413.11613954 532.08956315 377.91990213 534.80934537 448.04343602
595.27120684 666.70145306 511.81420024 572.88687124 504.99052047
565.20573493 460.22874533 449.68777644 422.84378902 456.61172298
597.67665595 449.88789011 615.14845865 511.67758056 504.18999612
515.72146994 568.16951591 551.48513998 356.6378314 464.8620116]
```

```

481.60170612 534.06829437 256.58748272 505.21497329 519.95716322
315.21654441 502.09614923 387.31041674 473.06173507 432.80626366
539.7086216 589.64106977 751.94849817 558.06069456 523.59234931
431.86945324 425.51361371 518.75816581 641.45267668 481.82884822
549.43958683 381.11034334 555.11914644 403.48664165 472.59054176
501.74908395 473.46525834 456.52751619 554.4243961 702.50615043
534.72434807 618.85560624 499.94890064 559.24813999 574.63905469
505.12302239 529.6719221 479.20777206 424.80282544 452.15936464
525.51286607 556.47232592 425.61830232 588.52893723 490.82444018
562.38936774 495.61462789 445.50430092 456.58675234 537.97742609
367.01103033 421.30159286 551.35810069 527.96988425 493.41227515
495.00368427 519.64738812 461.03646636 528.67738089 442.80460392
543.01475135 350.27066003 401.38884281 606.52218234 576.70349109
524.12472044 553.92991615 507.89702237 505.40680844 371.67040153
342.81954951 633.98898262 523.5392842 532.54259043 574.29573967
435.50713235 599.56950382 487.08846216 457.54189256 425.06792835
332.0253581 443.73046773 563.27115892 466.20350336 463.54617144
381.44663484 411.93032149 473.50490766 573.08604675 417.67287092
543.38703479 547.58495414 547.56984661 450.94665355 561.1959677
478.3032658 484.22533552 457.55220485 411.74640875 375.68250566
449.40702218 557.63342169 553.14582142 485.92999788 547.94903123
543.1115209 451.86394232 532.29902785 548.49324648 445.98203326
429.88352809 486.61675893 609.72814219 569.93904011 441.12529378]

```

```



In [7]: # To know how far are we from the actual price compared to predictions using the
# we us can scatter plot from matplotlib with y_test which never used to train a
plt.xlabel('Y Test (True Values)')
plt.ylabel('Y label Predicted Value')
print(plt.scatter(y_test, predictions))

# Print the Histogram of residuals; I see it is normally distributed which is go
print(sns.displot((y_test - predictions)))

# The three common Regression evaluation metrics are Mean Absolute Error(MAE), M
# The Lower we minimize the error the best the model
print('MAE')
print(
    metrics.mean_absolute_error(y_test, predictions)
)
print('MSE')
print(metrics.mean_squared_error(y_test, predictions))
print('Root MSE')
print(np.sqrt(metrics.mean_squared_error(y_test, predictions)))

<matplotlib.collections.PathCollection object at 0x00000279905C5160>
<seaborn.axisgrid.FacetGrid object at 0x00000279905C6CC0>
MAE
7.294546588331314
MSE
81.90726984520933
Root MSE
9.050263523522911


```



```
In [15]: print('Residuals')
sns.displot(y_test-predictions, bins=50)
```

Residuals

Out[15]: <seaborn.axisgrid.FacetGrid at 0x2799162f5c0>



In []: