

```
In [51]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
import seaborn as sns
from sklearn import metrics
%matplotlib inline
```

```
In [11]: print("Titanic survival prediction Using Logistic Regression Model from scikit-learn...")
print("☕ Grab a cup of coffee and watch the magic happen!")

# Load the dataset
train = pd.read_csv('../data/titanic_train.csv')
train.head
```


Titanic survival prediction Using Logistic Regression Model from scikit-learn...  
☕ Grab a cup of coffee and watch the magic happen!

```
Out[11]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.25
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.28
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.92
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.10
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.05

```
In [66]: #To check if any datapoints are null ; the particular chart below shows age data
sns.heatmap(train.isnull(), yticklabels=False, cbar=False, cmap='viridis')
```


```
Out[66]: <Axes: >
```



```
In [19]: sns.set_style('whitegrid')
```


```
In [70]: sns.countplot(x='Survived',hue='Pclass', data=train)
```

```
Out[70]: <Axes: xlabel='Survived', ylabel='count'>
```




```
In [74]: sns.displot(train['Age'].dropna(), kde=False, bins=30)
```

```
Out[74]: <seaborn.axisgrid.FacetGrid at 0x217308d7bf0>
```



```
In [35]: train['Age'].plot.hist(bins=35)
```

```
Out[35]: <Axes: ylabel='Frequency'>
```




```
In [37]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   PassengerId 891 non-null    int64
 1   Survived    891 non-null    int64
 2   Pclass      891 non-null    int64
 3   Name        891 non-null    object
 4   Sex         891 non-null    object
 5   Age         714 non-null    float64
 6   SibSp       891 non-null    int64
 7   Parch       891 non-null    int64
 8   Ticket      891 non-null    object
 9   Fare        891 non-null    float64
10   Cabin       204 non-null    object
11   Embarked    889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```


```
In [39]: sns.countplot(x='SibSp', data=train)
```

```
Out[39]: <Axes: xlabel='SibSp', ylabel='count'>
```




```
In [43]: train['Fare'].hist(bins=40, figsize=(10,4))
```

```
Out[43]: <Axes: >
```



```
In [78]: sns.boxplot(x='Pclass', y='Age', data=train)
```

```
Out[78]: <Axes: xlabel='Pclass', ylabel='Age'>
```




```
In [98]: def impute_age(cols):
Age = cols.iloc[0]
Pclass = cols.iloc[1]
if pd.isnull(Age): # Check if Age is NaN
    if Pclass == 1:
        return 37 # Impute age for Pclass 1
    elif Pclass == 2:
        return 29 # Impute age for Pclass 2
    else:
        return 24 # Impute age for Pclass 3 or any other condition
else:
    return Age # Return the original Age if it's not NaN
```

```
In [102... train['Age'] = train[['Age', 'Pclass']].apply(impute_age,axis=1)
```

```
In [104... sns.heatmap(train.isnull(), yticklabels=False, cbar=False, cmap = 'viridis')
```

```
Out[104... <Axes: >
```



```
In [94]: train.drop('Cabin', axis=1, inplace=True)
```

```
In [132... sex = pd.get_dummies(train['Sex'], drop_first=True)
```

```

-----
KeyError                                Traceback (most recent call last)
File C:\learnings\Lib\site-packages\pandas\core\indexes\base.py:3805, in Index.get_loc(self, key)
    3804 try:
-> 3805     return self._engine.get_loc(casted_key)
    3806 except KeyError as err:

File index.pyx:167, in pandas._libs.index.IndexEngine.get_loc()

File index.pyx:196, in pandas._libs.index.IndexEngine.get_loc()

File pandas\_libs\hashtable_class_helper.pxi:7081, in pandas._libs.hashtable.PyObjectHashTable.get_item()

File pandas\_libs\hashtable_class_helper.pxi:7089, in pandas._libs.hashtable.PyObjectHashTable.get_item()

KeyError: 'Sex'

The above exception was the direct cause of the following exception:

KeyError                                Traceback (most recent call last)
Cell In[132], line 1
----> 1 sex = pd.get_dummies(train['Sex'], drop_first=True)

File C:\learnings\Lib\site-packages\pandas\core\frame.py:4102, in DataFrame.__getitem__(self, key)
    4100 if self.columns.nlevels > 1:
    4101     return self._getitem_multilevel(key)
-> 4102 indexer = self.columns.get_loc(key)
    4103 if is_integer(indexer):
    4104     indexer = [indexer]

File C:\learnings\Lib\site-packages\pandas\core\indexes\base.py:3812, in Index.get_loc(self, key)
    3807 if isinstance(casted_key, slice) or (
    3808     isinstance(casted_key, abc.Iterable)
    3809     and any(isinstance(x, slice) for x in casted_key)
    3810 ):
    3811     raise InvalidIndexError(key)
-> 3812     raise KeyError(key) from err
    3813 except TypeError:
    3814     # If we have a listlike key, _check_indexing_error will raise
    3815     # InvalidIndexError. Otherwise we fall through and re-raise
    3816     # the TypeError.
    3817     self._check_indexing_error(key)

KeyError: 'Sex'

```

```
In [118... embark = pd.get_dummies(train['Embarked'], drop_first=True)
```

```
In [120... train = pd.concat([train,sex, embark], axis=1)
```

```
In [122... train.head()
```



Out[122...

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	24.0	1	0	A/5 21171	7.2
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	37.0	1	0	PC 17599	71.2
2	3	1	3	Heikkinen, Miss. Laina	female	24.0	0	0	STON/O2. 3101282	7.9
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	37.0	1	0	113803	53.1
4	5	0	3	Allen, Mr. William Henry	male	24.0	0	0	373450	8.0



In [134...

```
train.drop(['Sex', 'Embarked', 'Name', 'Ticket'], axis=1, inplace=True)
```

```

-----
KeyError                                Traceback (most recent call last)
Cell In[134], line 1
----> 1 train.drop(['Sex', 'Embarked', 'Name', 'Ticket'], axis=1, inplace=True)

File C:\learnings\Lib\site-packages\pandas\core\frame.py:5581, in DataFrame.drop
(self, labels, axis, index, columns, level, inplace, errors)
    5433 def drop(
    5434     self,
    5435     labels: IndexLabel | None = None,
    (... )
    5442     errors: IgnoreRaise = "raise",
    5443 ) -> DataFrame | None:
    5444     """
    5445     Drop specified labels from rows or columns.
    5446
    (... )
    5579         weight  1.0      0.8
    5580     """
-> 5581     return super().drop(
    5582         labels=labels,
    5583         axis=axis,
    5584         index=index,
    5585         columns=columns,
    5586         level=level,
    5587         inplace=inplace,
    5588         errors=errors,
    5589     )

File C:\learnings\Lib\site-packages\pandas\core\generic.py:4788, in NDFrame.drop
(self, labels, axis, index, columns, level, inplace, errors)
    4786 for axis, labels in axes.items():
    4787     if labels is not None:
-> 4788         obj = obj._drop_axis(labels, axis, level=level, errors=errors)
    4790 if inplace:
    4791     self._update_inplace(obj)

File C:\learnings\Lib\site-packages\pandas\core\generic.py:4830, in NDFrame._drop
_axis(self, labels, axis, level, errors, only_slice)
    4828     new_axis = axis.drop(labels, level=level, errors=errors)
    4829     else:
-> 4830     new_axis = axis.drop(labels, errors=errors)
    4831     indexer = axis.get_indexer(new_axis)
    4833 # Case for non-unique axis
    4834 else:

File C:\learnings\Lib\site-packages\pandas\core\indexes\base.py:7070, in Index.dr
op(self, labels, errors)
    7068 if mask.any():
    7069     if errors != "ignore":
-> 7070         raise KeyError(f"{labels[mask].tolist()} not found in axis")
    7071     indexer = indexer[~mask]
    7072 return self.delete(indexer)

KeyError: "[ 'Sex', 'Embarked', 'Name', 'Ticket'] not found in axis"

```

```

In [136... train[['male', 'Q', 'S']] = train[['male', 'Q', 'S']].astype(int)
train.head()

```

Out[136...

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	1	0	3	24.0	1	0	7.2500	1	0	1
1	2	1	1	37.0	1	0	71.2833	0	0	0
2	3	1	3	24.0	0	0	7.9250	0	0	1
3	4	1	1	37.0	1	0	53.1000	0	0	1
4	5	0	3	24.0	0	0	8.0500	1	0	1

In [138...

```
train.drop('PassengerId', axis=1, inplace=True)
```

In [140...

```
train.head()
```

Out[140...

	Survived	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	0	3	24.0	1	0	7.2500	1	0	1
1	1	1	37.0	1	0	71.2833	0	0	0
2	1	3	24.0	0	0	7.9250	0	0	1
3	1	1	37.0	1	0	53.1000	0	0	1
4	0	3	24.0	0	0	8.0500	1	0	1

In [142...

```
x = train.drop('Survived', axis=1)
y = train['Survived']
```

In [144...

```
from sklearn.model_selection import train_test_split
```

In [146...

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random
```

In [148...

```
lgrm = LogisticRegression()
lgrm.fit(x_train, y_train)
```

C:\learnings\Lib\site-packages\sklearn\linear\_model\\_logistic.py:469: Convergence Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Out[148...

▼ LogisticRegression ⓘ ?

LogisticRegression()

In [150...

```
predictions = lgrm.predict(x_test)
```

In [152...

```
from sklearn.metrics import classification_report
```

In [154...

```
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.78	0.88	0.82	169
1	0.80	0.67	0.73	126
accuracy			0.79	295
macro avg	0.79	0.77	0.78	295
weighted avg	0.79	0.79	0.78	295

In [156... `from sklearn.metrics import confusion_matrix`

In [162... `cf = confusion_matrix(y_test, predictions)`

In [164... `cm_df = pd.DataFrame(cf,  
index=['Actual Negative', 'Actual Positive'],  
columns=['Predicted Negative', 'Predicted Positive'])`

In [168... `cm_df.head()`

Out[168...

	Predicted Negative	Predicted Positive
Actual Negative	148	21
Actual Positive	42	84

In [ ]: