

```
In [307... import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [359... #Load stock historical data
nvda_data = pd.read_csv('nvda.csv')
spy_data = pd.read_csv('spy.csv')
# Remove Volume column from both DataFrames
nvda_data.drop(columns=['Volume'], inplace=True)
spy_data.drop(columns=['Volume'], inplace=True)
```

```
In [363... nvda_data.head()
```

```
Out[363...      Date  Close/Last  Open  High  Low
0  10/11/2024    $134.80  $134.01  $135.78  $133.66
1  10/10/2024    $134.81  $131.91  $135.00  $131.00
2  10/09/2024    $132.65  $134.11  $134.52  $131.38
3  10/08/2024    $132.89  $130.26  $133.48  $129.42
4  10/07/2024    $127.72  $124.99  $130.64  $124.95
```

```
In [365... spy_data.head()
```

```
Out[365...      Date  Close/Last  Open  High  Low
0  10/11/2024    579.58  576.05  580.3300  575.9100
1  10/10/2024    576.13  575.77  577.5800  574.4900
2  10/9/2024    577.14  573.16  577.7100  572.5500
3  10/8/2024    573.17  570.42  573.7800  569.5299
4  10/7/2024    567.80  571.30  571.9599  566.6300
```

```
In [373... nvda_data['Date'] = pd.to_datetime(nvda_data['Date'])
spy_data['Date'] = pd.to_datetime(spy_data['Date'])
combined_data = pd.merge(nvda_data, spy_data, on='Date', suffixes=('_nvda', '_sp
combined_data.head()
```

Out[373...

	Date	Close/Last_nvda	Open_nvda	High_nvda	Low_nvda	Close/Last_spy	Open_spy
0	2024-10-11	\$134.80	\$134.01	\$135.78	\$133.66	579.58	576.05
1	2024-10-10	\$134.81	\$131.91	\$135.00	\$131.00	576.13	575.77
2	2024-10-09	\$132.65	\$134.11	\$134.52	\$131.38	577.14	573.16
3	2024-10-08	\$132.89	\$130.26	\$133.48	\$129.42	573.17	570.42
4	2024-10-07	\$127.72	\$124.99	\$130.64	\$124.95	567.80	571.30



In [375...

```

# Check for NaN values in the specified columns
import pandas as pd
import matplotlib.pyplot as plt

# Assuming combined_data is already defined and contains the necessary data

# Create a mask for rows with NaN values in the specified columns
nan_mask = combined_data[['Open_nvda', 'Close/Last_nvda']].isna().any(axis=1)

# Extract dates with NaN values
nan_dates = combined_data[nan_mask]['Date']

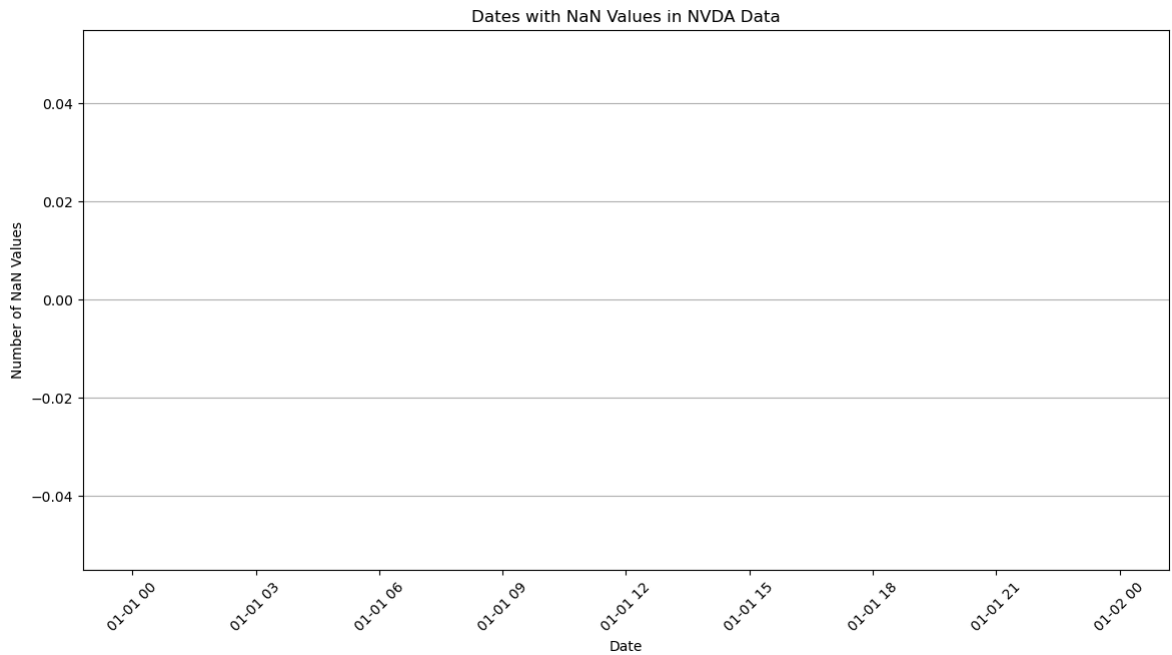
# Count the occurrences of NaN values per date
nan_counts = nan_dates.value_counts().sort_index()

# Plotting the results
plt.figure(figsize=(14, 7))
plt.bar(nan_counts.index, nan_counts.values, color='red', alpha=0.7)
plt.title('Dates with NaN Values in NVDA Data')
plt.xlabel('Date')
plt.ylabel('Number of NaN Values')
plt.xticks(rotation=45)
plt.grid(axis='y')

plt.show()

print("NaN values in relevant columns:")
print(combined_data[['Open_nvda', 'Close/Last_nvda', 'Open_spy', 'Close/Last_spy']])

```



NaN values in relevant columns:

```
Open_nvda      0
Close/Last_nvda 0
Open_spy       0
Close/Last_spy 0
dtype: int64
```

```
In [377... # Check for NaN values in the specified columns
print("NaN values in relevant columns:")
print(combined_data[['Open_nvda', 'Close/Last_nvda', 'Open_spy', 'Close/Last_spy']])
```

NaN values in relevant columns:

```
Open_nvda      0
Close/Last_nvda 0
Open_spy       0
Close/Last_spy 0
dtype: int64
```

```
In [379... # Define columns to clean and convert
cols_to_clean = ['Open_nvda', 'Close/Last_nvda', 'Open_spy', 'Close/Last_spy']

# Convert columns to string to handle non-string types
for col in cols_to_clean:
    combined_data[col] = combined_data[col].astype(str) # Convert to string

# Clean the columns by removing dollar signs and extra spaces, then convert to numeric
for col in cols_to_clean:
    combined_data[col] = combined_data[col].str.replace('$', '', regex=False) # Remove dollar signs
    combined_data[col] = combined_data[col].str.strip() # Remove any leading/trailing spaces
    combined_data[col] = pd.to_numeric(combined_data[col], errors='coerce') # Convert to numeric
```

```
In [381... cols_to_check = ['Open_nvda', 'Close/Last_nvda', 'Open_spy', 'Close/Last_spy']
scaler = StandardScaler()
combined_data[cols_to_check] = scaler.fit_transform(combined_data[cols_to_check])
print("Scaled data:")
print(combined_data.head())
```

Scaled data:

	Date	Close/Last_nvda	Open_nvda	High_nvda	Low_nvda	Close/Last_spy	\
0	2024-10-11	4.317199	4.287475	\$135.78	\$133.66	2.468255	
1	2024-10-10	4.317568	4.210020	\$135.00	\$131.00	2.434612	
2	2024-10-09	4.237884	4.291164	\$134.52	\$131.38	2.444461	
3	2024-10-08	4.246738	4.149162	\$133.48	\$129.42	2.405747	
4	2024-10-07	4.056012	3.954784	\$130.64	\$124.95	2.353382	

	Open_spy	High_spy	Low_spy
0	2.434830	580.3300	575.9100
1	2.432099	577.5800	574.4900
2	2.406643	577.7100	572.5500
3	2.379919	573.7800	569.5299
4	2.388502	571.9599	566.6300

```
In [383... # Moving Averages
combined_data['MA50_nvda'] = combined_data['Close/Last_nvda'].rolling(window=50)
combined_data['MA200_nvda'] = combined_data['Close/Last_nvda'].rolling(window=200)
combined_data['MA50_spy'] = combined_data['Close/Last_spy'].rolling(window=50).r
combined_data['MA200_spy'] = combined_data['Close/Last_spy'].rolling(window=200)

# Daily Returns
combined_data['Returns_nvda'] = combined_data['Close/Last_nvda'].pct_change()
combined_data['Returns_spy'] = combined_data['Close/Last_spy'].pct_change()

# Volatility (rolling 30-day standard deviation)
combined_data['Volatility_spgi'] = combined_data['Returns_nvda'].rolling(window=
combined_data['Volatility_spy'] = combined_data['Returns_spy'].rolling(window=30

# Drop NaN values that result from rolling calculations
combined_data.dropna(inplace=True)
```

```
In [385... # Label 1 if SPGI underperforms SPY, otherwise 0
combined_data['Underperform'] = (combined_data['Returns_nvda'] < combined_data['

# Optional: Shift the target label by 1 day if you're predicting the next day's
combined_data['Target'] = combined_data['Underperform'].shift(-1)

# Drop the last row because of the shift
combined_data.dropna(inplace=True)
```

```
In [387... from sklearn.model_selection import train_test_split

# Select features for training
features = ['Open_nvda', 'Close/Last_nvda', 'MA50_nvda', 'MA200_nvda', 'Returns
            'Open_spy', 'Close/Last_spy', 'MA50_spy', 'MA200_spy', 'Returns_spy

X = combined_data[features] # Features
y = combined_data['Target'] # Target: Whether SPGI underperforms SPY

# Train-test split (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle
```

```
In [391... from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Initialize and train the model
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```

```
# Predict on the test set
y_pred = model.predict(X_test)

# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")

# Detailed classification report
print(classification_report(y_test, y_pred))
```

Model Accuracy: 0.50

	precision	recall	f1-score	support
0.0	0.49	0.99	0.66	226
1.0	0.78	0.03	0.06	238
accuracy			0.50	464
macro avg	0.64	0.51	0.36	464
weighted avg	0.64	0.50	0.35	464


In [393...

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load your combined_data DataFrame, which should contain SPGI and SPY data
# combined_data should already have Date, Close/Last_spgi, Close/Last_spy column

# Plotting
plt.figure(figsize=(14, 7))
plt.plot(combined_data['Date'], combined_data['Close/Last_nvda'], label='NVDA',
plt.plot(combined_data['Date'], combined_data['Close/Last_spy'], label='SPY', co

plt.title('NVDA vs. SPY Closing Prices')
plt.xlabel('Date')
plt.ylabel('Price')
plt.legend()
plt.grid()
plt.show()
```



In [395...

```


# Calculate daily returns
combined_data['Return_spgi'] = combined_data['Close/Last_nvda'].pct_change()
combined_data['Return_spy'] = combined_data['Close/Last_spy'].pct_change()

# Drop NaN values for returns
combined_data.dropna(inplace=True)

# Plotting returns
plt.figure(figsize=(14, 7))
plt.plot(combined_data['Date'], combined_data['Return_spgi'], label='NVDA Return')
plt.plot(combined_data['Date'], combined_data['Return_spy'], label='SPY Returns')

plt.title('NVDA vs. SPY Daily Returns')
plt.xlabel('Date')
plt.ylabel('Daily Return')
plt.legend()
plt.grid()
plt.show()

```



In [398...

```

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
combined_data[cols_to_check] = scaler.fit_transform(combined_data[cols_to_check])

```

In [400...

```

from sklearn.preprocessing import RobustScaler

scaler = RobustScaler()
combined_data[cols_to_check] = scaler.fit_transform(combined_data[cols_to_check])

```

In [402...

```

from sklearn.preprocessing import PowerTransformer

scaler = PowerTransformer()
combined_data[cols_to_check] = scaler.fit_transform(combined_data[cols_to_check])

```

In [ ]:

```

from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],

```

```
    'min_samples_split': [2, 5, 10]
}

grid_search = GridSearchCV(RandomForestClassifier(random_state=42), param_grid,
grid_search.fit(X_train, y_train)

print("Best parameters:", grid_search.best_params_)
```

```
In [ ]: from sklearn.ensemble import GradientBoostingClassifier
```

```
model = GradientBoostingClassifier(random_state=42)
model.fit(X_train, y_train)
```

```
In [ ]: # Predict on the test set
y_pred = model.predict(X_test)

# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")

# Detailed classification report
print(classification_report(y_test, y_pred))
```

```
In [ ]:
```